

March 9, 2026

Editor-in-Chief
Data in Brief
Elsevier

Dear Editor,

I am submitting a data article entitled “**SALAD: SOC Alert Labeled Analysis Dataset — A Unified Benchmark for Alert Triage, Prioritization, and Classification**” for consideration in *Data in Brief*.

Security Operations Center (SOC) alert triage is a pressing challenge: analysts face over 4,000 alerts daily, yet no public dataset provides SOC alerts with triage decision labels, priority scores, and MITRE ATT&CK mapping in a unified, benchmark-ready format. Existing public datasets such as CIC-IDS2017 and UNSW-NB15 offer raw network traffic without the SOC-specific metadata that researchers in machine learning, reinforcement learning, and large language models need to develop and compare alert triage systems. SALAD fills this gap. It contains 2,778,424 alerts in a standardized 29-field schema with five layers of computed metadata—triage decisions, priority scores, difficulty levels, MITRE ATT&CK labels, and natural language descriptions—none of which exist in the source data. The dataset defines four benchmark tasks with baseline results and evaluation scripts, and it is publicly available on Zenodo (DOI: 10.5281/zenodo.18912306), HuggingFace, and GitHub under a CC-BY 4.0 license.

This manuscript is a standalone data article. It has not been published elsewhere and is not under consideration at any other journal. I declare no competing interests.

Sincerely,

Nutthakorn Chalaemwongwan
Department of Computer
Engineering, KOSEN-KMITL
King Mongkut’s Institute of
Technology Ladkrabang
Bangkok, Thailand
nutthakorn.ch@kmitl.ac.th